

When compact latents collapse: a retained negative ARC evaluation

Ryan Gomez | Author review draft | Retained evidence

RESEARCH QUESTION

Does adding quantization, memory and a latent planner produce useful representation compression under a frozen comparison?

METHOD AND EVIDENCE BOUNDARY

Hashed token embeddings, mean pooling, a 32-code quantizer, sparse memory and one latent action step; same input target alignment. This is not a Transformer or canonical context to distinct target prediction.

Five fixed seeds; 1,117 training and 295 validation items; gradient active parameters matched within 1%. Report arithmetic is rechecked here; no training or checkpoint replay.


WHY THIS MATTERS FOR PARSIMONY

Compression and collapse are different. Inspect code use, input dependence, task accuracy and positive controls together. A mechanism ablation inside a collapsed full model cannot establish a general verdict on planning or target alignment.

WHAT REMAINS UNRESOLVED

A later trainability repair remained negative or inconclusive. The protected confirmatory test stays unopened for this failed hypothesis.

Frozen validation accuracy



Mean +/- sample SD across five fixed seeds. 295 validation items per seed.

WHAT THE RESULT SUPPORTS

Full minus matched accuracy: -0.0115, below the declared +0.02 practical gain. The superiority gate failed. Planner and target attribution gates also failed.

Retained review describes single-code quantizer use and constant predictions. Disabling quantization restored input dependence without establishing above chance accuracy.